



I



II



III

frames · audio · text

(a) Clips I-III

faces → ArcFace
→ identities



most frames
in clip III



2nd most
in clip III



all others

no clip IV, no names



(b) Role assignment

speech

scene

face token = attention pooling over
the role's frames + role/clip embedding

(c) Evidence set (15 tokens)



(d) Set encoder

$p(y_B)$, clip IV

ang

dis

fea

hap

neu

sad

sur