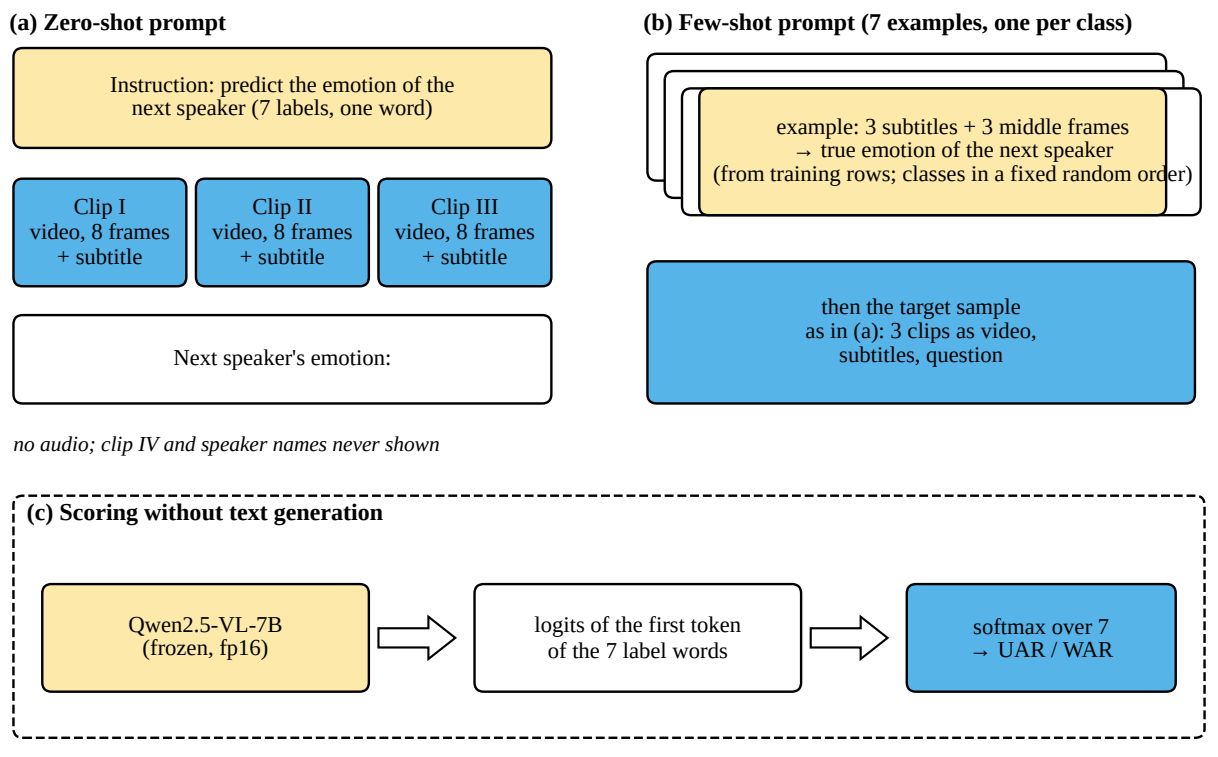
(a) Zero-shot prompt
Instruction: predict the emotion of the next speaker (7 labels, one word)
Clip I video, 8 frames + subtitle
Clip II video, 8 frames + subtitle
Clip III video, 8 frames + subtitle
Next speaker's emotion:
no audio; clip IV and speaker names never shown
(b) Few-shot prompt (7 examples, one per class)
example: 3 subtitles + 3 middle frames → true emotion of the next speaker (from training rows; classes in a fixed random order)
then the target sample as in (a): 3 clips as video, subtitles, question
(c) Scoring without text generation
Qwen2.5-VL-7B (frozen, fp16)
logits of the first token of the 7 label words
softmax over 7 → UAR / WAR