



Clip I (context)



Clip II: speaker A



Clip III: listener L

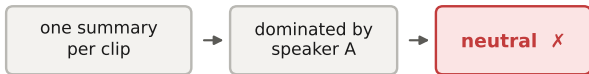


Clip IV: responder B (unseen)

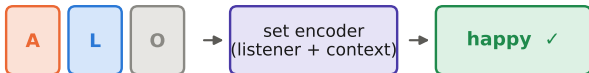


Who will answer is already on screen, listening.

Speaker-centric (previous work)



Listener-centric (RoleNet)



person x turn tokens