




Clips I-III



clip I: context



clip III: speaker A

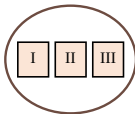


clip III: listener L

clip IV (not an input):
 $B = L$ answers

Forecast for B

previous methods



*clip summaries
(speaker dominates)*

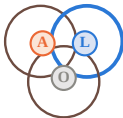
happy

neutral X

sad

Forecast for B

our method



person $\times$ turn tokens

happy ✓

neutral

sad